

# Futureverse P2P: Peer-to-Peer Parallelization in R

- *Share compute among friends across the world*



**Henrik Bengtsson**

University of California, San Francisco

R Foundation, R Consortium

 @HenrikBengtsson

useR! 2025, Durham, NC, USA (2025-08-10)



# Future ... what?

An R assignment:

```
a <- 1 + 2
```

```
x <- 1:10  
b <- slow_sum(x)
```

A future-value assignment:

```
f <- future({ 1 + 2 })  
a <- value(f)
```

```
x <- 1:10  
f <- future({ slow_sum(x) })  
b <- value(f)
```



# Future ... why? (bc enables parallel processing)

Two calculations:

```
1 x_a <- 1:10  
2 x_b <- 11:20  
3 a <- slow_sum(x_a) # 1 min  
4 b <- slow_sum(x_b) # 1 min  
5 c <- a + b
```

Total time: 2 mins

Two futures:

```
1 x_a <- 1:10  
2 x_b <- 11:20  
3 f_a <- future( slow_sum(x_a) ) # 0 sec  
4 f_b <- future( slow_sum(x_b) ) # 0 sec  
5  
6 a <- value(f_a) # 1 min  
7 b <- value(f_b)  
8 c <- a + b
```

Total time: 1 min

=> futures are the core building block for parallel processing

# Futureverse allows you to stick with your favorite coding style

Parallel alternatives to traditional, sequential functions:

```
1 ys <- lapply(xs, slow_sum)                                # base R
2 ys <- future_lapply(xs, slow_sum)                            # {future.apply}
```

```
1 ys <- map(xs, slow_sum)                                    # {purrr}
2 ys <- future_map(xs, slow_sum)                             # {furrr}
```

```
1 plan(multisession)
2 plan(future.callr::callr)
3 plan(future.mirai::mirai_multisession)
4 plan(cluster, workers = c("n1", "desktop", "server.myuniv.org"))
5 plan(future.batchtools::batchtools_slurm)
```

Let's go back in time...

# Ten-Year Anniversary (future 0.6.0 released June 2015)

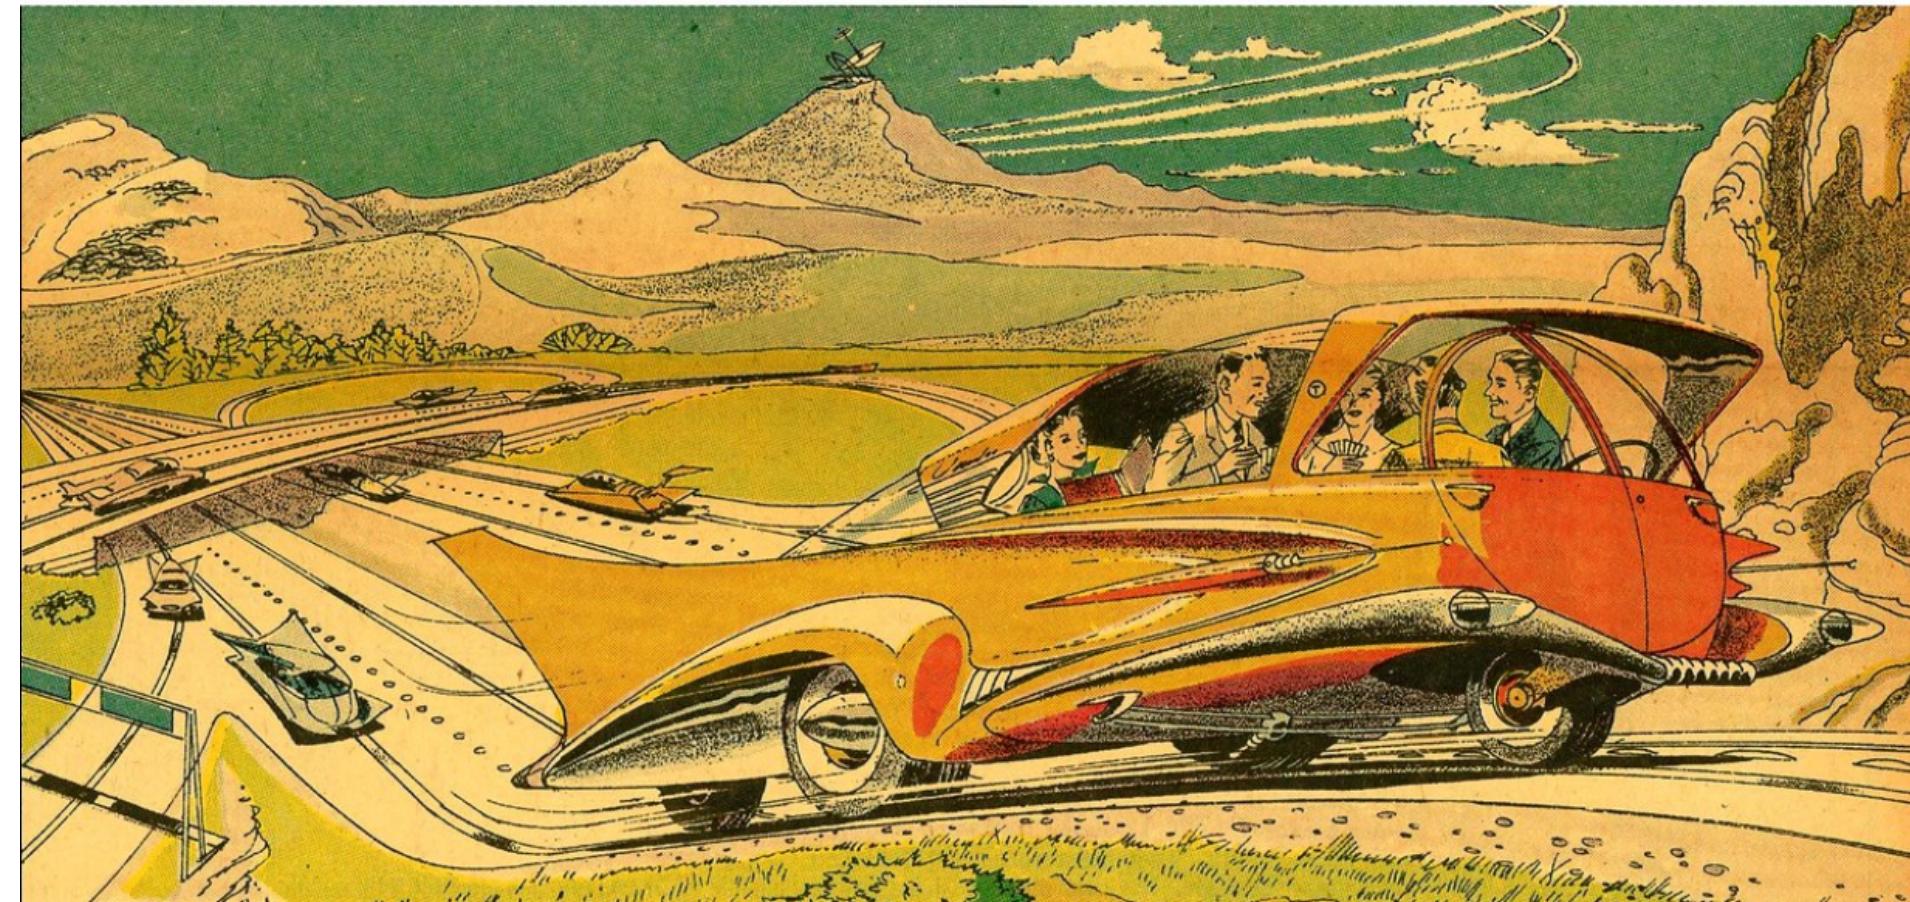


future logo: *Dan LaBar*, ggplot2 logo balloon wall: *Greg Swinehart & Hadley Wickham*

# useR! 2016 at Stanford, CA, USA

## A Future for R

Henrik Bengtsson, UC San Francisco



useR 2016, Stanford, CA, 2016-06-28

# useR! 2016 at Stanford, CA, USA

## R package: future



- A simple unified API
- Works the same on all platforms
- Easy to install
- Lightweight (~300 kB incl. dependencies)
- Vignettes
- Extendable by anyone

CRAN 1.0.0 build passing Codecov 97%

5 / 18

# useR! 2016 at Stanford, CA, USA

## Consistent futures everywhere

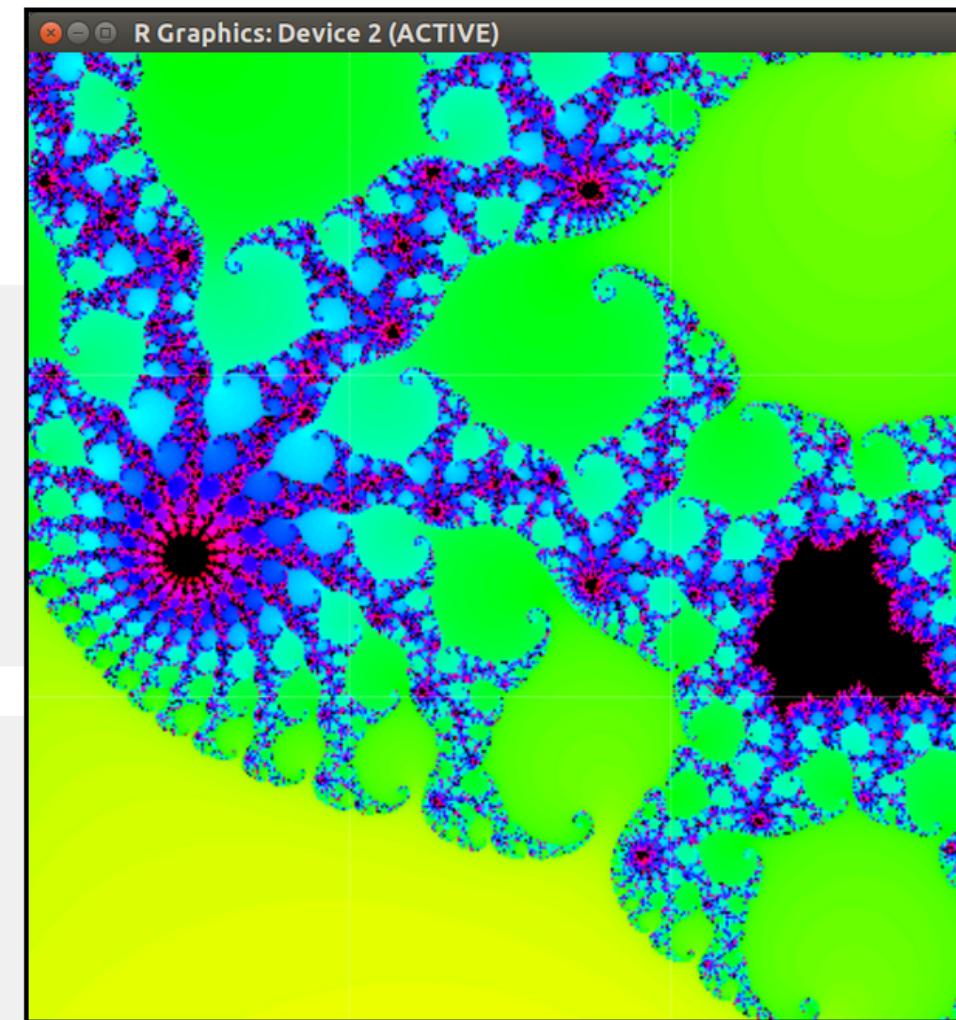


- Unix
- macOS
- Windows

```
> library("future")
> plan(multiprocess)
> demo("mandelbrot")
```

Calculating and plotting  
Mandelbrot regions ...

- Region 1 done
- Region 2 done
- Region 7 done
- Region 5 done
- ...



7 / 18

# useR! 2016 at Stanford, CA, USA

## future.BatchJobs: Futures for HPC

future.BatchJobs:	Job scheduler:
batchjobs_slurm	Slurm
batchjobs_sge	Sun Grid Engine
batchjobs_torque	TORQUE / PBS
batchjobs_lsf	Load Sharing Facility
batchjobs_openlava	OpenLava

```
library("future.BatchJobs")
plan(batchjobs_slurm)
```

```
bam <- listenv()
for (i in seq_along(fastq)) {
  bam[[i]] %<-% DNAseq::align(fastq[i])
}
```

CRAN 0.12.1 build passing Codecov 90%

12 / 18

# useR! 2016 at Stanford, CA, USA

## A7. Futures I'd like to see

- `plan(r32)`
  - e.g. `RODBC::odbcConnectAccess()` works only on 32-bit R.
- `plan(p2p)`
  - Private and / or community-based peer-to-peer computer cluster
- `plan(rhelp)`
  - Post R scripts to R-help and ask for the results :P



# Back to the present

# future.p2p: sharing compute among friends - it's easy!

```
library(future.apply)
```

```
plan(future.p2p::cluster, cluster = "alice/friends")
```

```
y <- future_lapply(xs, slow_sum)
```



# How to get started...

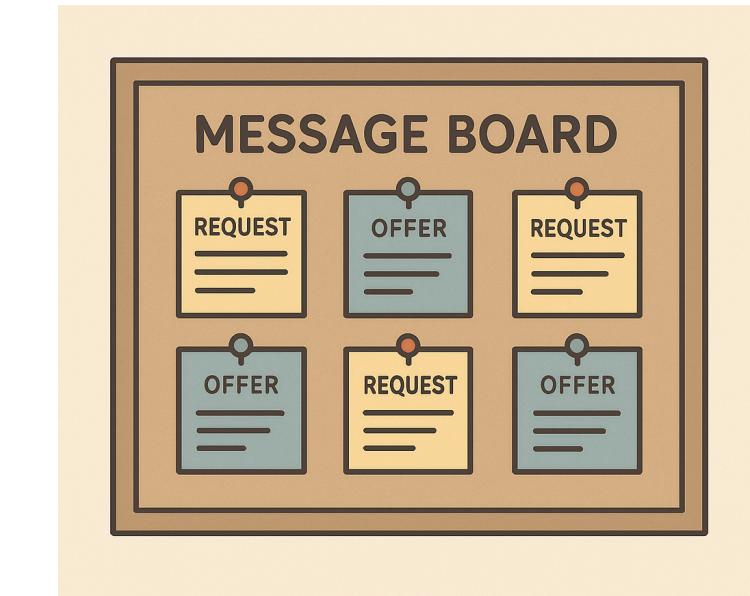
<https://future.p2p.futureverse.org>



# A P2P cluster has two components

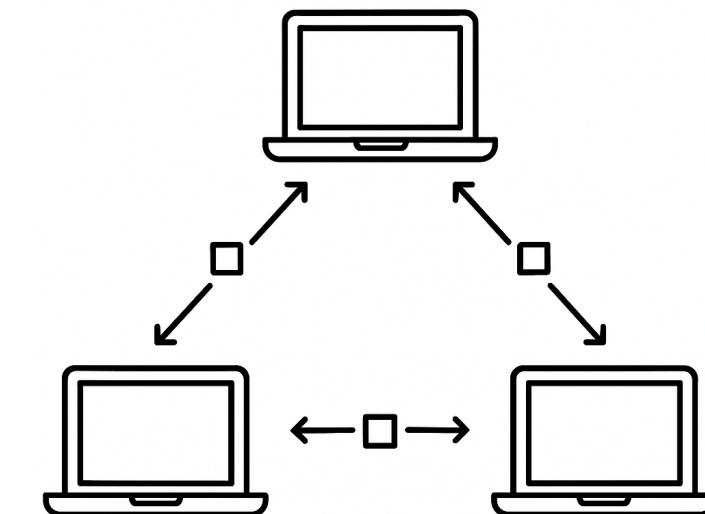
## Message Board

Used to announce futures and offers to do work  
(centralized; lightweight - only metadata)



## P2P file-transfer protocol

Used to send futures to workers and receive results  
(peer-to-peer; full-size data transfers)



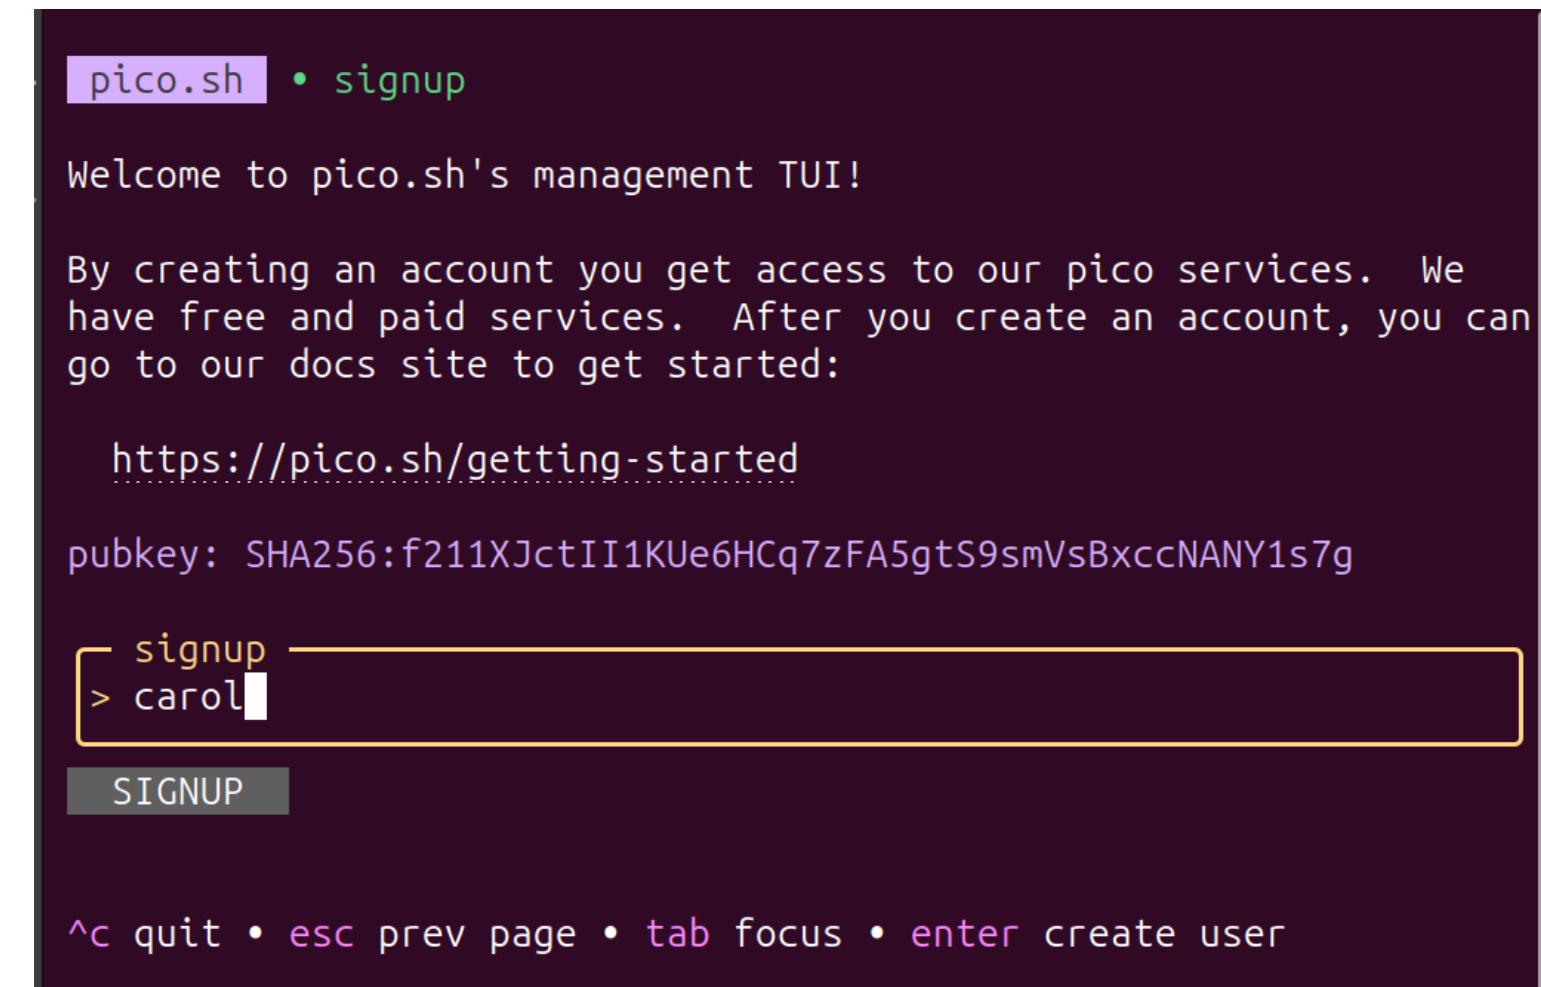
# To join a P2P cluster you need an account

All P2P cluster users need a [pico.sh](#) account to access the message board:

## 1. ssh-keygen

```
{carol}$ ssh-keygen
Generating public/private ed25519 key pair.
Enter file in which to save the key (/home/carol/.ssh/id_ed25519):
Created directory '/home/carol/.ssh'.
Enter passphrase (empty for no passphrase):
Enter same passphrase again:
Your identification has been saved in /home/carol/.ssh/id_ed25519
Your public key has been saved in /home/carol/.ssh/id_ed25519.pub
The key fingerprint is:
SHA256:f211XJctII1KUe6HCq7zFA5gtS9smVsBxccNANY1s7g carol@hb-x1-2025
The key's randomart image is:
---[ED25519 256]---
   ==0+*=oo. |
   o oo.+o... o|
   o . oo... ..+|
   . o + o... . oo|
     B E S o . +|
   . B o o . . . |
   . + . . . o |
   .o . . . |
   .|
```

## 2. ssh pico.sh => pick a username



That's it!

# Alice hosts a P2P cluster

Alice sets up P2P cluster and gives 'bob' and 'carol' access:

```
[alice]> future.p2p::host_cluster("alice/friends", users=c("bob", "carol"))
```

This is basically setting up a shared message board.

Q. What happens if Bob tries to use the P2P cluster?

```
[bob]> plan(future.p2p::cluster, cluster = "alice/friends")
[bob]> y <- future_lapply(xs, slow_sum)
```

Nothing - it will get stuck! Why?

# Contributing P2P workers is easy!

```
[alice]> future.p2p::worker("alice/friends")
```

```
[ bob ]> future.p2p::worker("alice/friends")
```

```
[carol]> future.p2p::worker("alice/friends")
```

```
[bob]> plan(future.p2p::cluster, cluster = "alice/friends")
[bob]> y <- future_lapply(xs, slow_sum)
```

=> Processed on three P2P workers

# You can set it up from the terminal

Setting up a P2P cluster:

```
{alice}$ Rscript -e future.p2p::host_cluster --cluster=alice/friends \
--users=bob,carol
```

Starting P2P workers:

```
{alice}$ Rscript -e future.p2p::worker --cluster=alice/friends &
{alice}$ Rscript -e future.p2p::worker --cluster=alice/friends &
```

```
{ bob }$ Rscript -e future.p2p::worker --cluster=alice/friends &
```

```
{carol}$ Rscript -e future.p2p::worker --cluster=alice/friends &
{carol}$ Rscript -e future.p2p::worker --cluster=alice/friends &
{carol}$ Rscript -e future.p2p::worker --cluster=alice/friends &
```

=> A shared P2P cluster with 6 workers. More can be added at any time!

# Pros and cons



# High latency but also high throughput

- High latency:
  - round-trip takes time, because p2p file transfers take time
  - Example:  $1+2$  takes 2-10 seconds to send, evaluate, and return
  - Just a prototype, so this will be improved
- High throughput:
  - any number of users and workers can join
  - a public P2P cluster could have thousands of P2P workers

# ⚠️ P2P computing requires mutual trust ⚠️

What if?

```
f <- future(system("erase-harddrive"))
```

Some technical mitigations:

- End-to-end encryption
- Run workers in sandboxed environments, e.g.
  - in a Linux container, or
  - in a virtual machine

See **parallelly** package for some examples

=> I encourage work on these topics - it's challenging, but important

The Future is bright  
... and fabulous

# Stay tuned for exciting Futureverse improvements

## On the horizon

- Custom random number generators (RNG) [together with Ralf Stubner]
- Resource specifications, e.g.
  - avoid memory overuse, e.g. `memory=2*GiB`
  - distribute to proper machines, e.g. `memory=2*GiB` and `gpu`

## In the near future

- If you think `furrr` and `future.apply` are neat - just wait!

# Thank you and may the future be with you!

- It's easy to get started - just try it
- Support: <https://github.com/HenrikBengtsson/future/discussions>
- Tutorials: <https://www.futureverse.org/tutorials.html>
- Blog posts: <https://www.futureverse.org/blog.html>
- More features on the roadmap
- I love feedback, ideas, and bug reports ❤️

